

Overview:

- Modified the scheduler to use a round-robin system.
- Added support for priority and priority aging.
- Added support for performance-metric tracking (bonus part 3)

Infrastructure Changes

```
int wait(int* status); // Lab1: updated
int waitpid(int pid, int *status, int options); // Lab1: Added to enable waitpid
int setpriority(int priority); // Lab2: Added to enable setpriority
```

defs.h: necessary system-call infrastructure for setpriority

```
// Per-process state
struct proc {
    uint sz; // Size of process memory (bytes)
    pde_t* pgdir; // Page table
    char *kstack; // Bottom of kernel stack for this process
    enum procstate state; // Process state
    int pid; // Process ID
    struct proc *parent; // Parent process
    struct trapframe *tf; // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan; // If non-zero, sleeping on chan
    int killed; // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd; // Current directory
    char name[16]; // Process name (debugging)
    int exitstatus; // Process exit status

    int priority; // Lab 2: The priority value. Range is 0-31

    int starttime; // Lab 2: Tracks start time (bonus points)
    int waittime; // Lab 2: Tracks waited time (bonus points)
};
```

proc.h: proc struct now has fields for tracking performance metrics and priority

```
np->starttime = ticks; // Lab 2: Updates start time of new process
```

Proc.c, fork(): forked procs store start time

```
int elapsedtime = ticks - curproc->starttime; // Lab2: Calculate elapsed time since start
cprintf("Turnaround (elapsed) time: %d %s\n", elapsedtime, "ticks."); // Lab 2: Print (bonus)
cprintf("Wait time: %d %s\n", curproc->waittime, "ticks."); // Lab 2: Print (bonus)
```

proc.c, exit(): exiting procs report their performance metrics

```
// Lab 2: Sets the priority of a process. This is a system call.
int setpriority(int priority) {

    struct proc *curproc = myproc(); // Grab the current process

    if (!curproc) { // Check if we failed to grab.
        return -1; // Terminate
    }

    acquire(&ptable.lock); // Get the semaphore to modify the priority of the current process
    curproc->priority = priority; // Modify the priority of the current process
    release(&ptable.lock); // Release the lock

    return 0;
}
```

proc.c: Definition of the setpriority system calls

```
extern int sys_setpriority(void); // Lab 3: Added to enable setpriority
```

syscall.c: necessary system call infrastructure

```
#define SYS_setpriority 5 // Lab 2: Added to enable setpriority
```

syscall.h: necessary system-call infrastructure

```
// Lab 2: Created to enable setpriority
int
sys_setpriority(void)
{
    char* priority;
    if(argptr(0, &priority, 8) < 0) return -1;

    return setpriority((int)priority);
}
```

sysproc.c: System call handler for setpriority

```

struct proc *s;
// Lab 1: Increase the priority of all the waiting processes. LOWER IS BETTER
for (s = ptable.proc; s < &ptable.proc[NPROC]; s++) {
    if (s -> priority > 0) s -> priority = s -> priority - 1;
}
// Lab 2: Set p to the process with the highest priority (the lowest value) that is also ready
for (s = ptable.proc; s < &ptable.proc[NPROC]; s++) {
    if (s -> state == RUNNABLE && s -> priority < p -> priority) p = s;
}

for (s = ptable.proc; s < &ptable.proc[NPROC]; s++) {
    if (s->state != RUNNABLE) continue;

    if (s != p) s -> waittime = ticks - s -> starttime; // Ensure that we aren't changing the currently-

}

// Switch to chosen process. It is the process's job
// to release ptable.lock and then reacquire it
// before jumping back to us.
c->proc = p;
switchvm(p);
p->state = RUNNING;

swtch(&(c->scheduler), p->context);
switchkvm();

if (s -> priority <= 29) {
    p -> priority += 2; // Lab 2: Decrease the priority the running process.
                       // By 2 since only doing 1 would cause it to
                       // slingshot back and forth, defeating the purpose of aging.
} else if (s->priority < 31) {
    p -> priority += 1;
}
}

```

proc.c, scheduler(): 3 for loops pass over the proc table to raise the priority of the waiting procs, set the current proc to the highest-priority waiting proc, and update the wait-times of all waiting procs. This is done in order. Finally, the priority of the running proc is reduced.

```
int setpriority(int); // Lab 1: Added declaration
```

User.h: necessary system-call infrastructure

```
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ lab2 2

This program tests the correctness of your lab#2

Part c) testing that wait and waitpid interact correctly with the roundrobin scheduler:

The is child with PID# 4
Turnaround (elapsed) time: 1 ticks.
Wait time: 0 ticks.

The is child with PID# 5
Turnaround (elapsed) time: 2 ticks.
Wait time: 1 ticks.

The is child with PID# 6
Turnaround (elapsed) time: 0 ticks.

The is child with PID# 7
Turnaround (elapsed) time: 1 ticks.
Wait time: 0 ticks.

The is child with PID# 8
Turnaround (elapsed) time: 2 ticks.
Wait time: 1 ticks.
Wait time: 2 ticks.

This is the parent: Now waiting for child with PID# 7
This is the parent: Child# 7 has exited
This is the parent: Now waiting for child with PID# 5
This is the parent: Child# 5 has exited
This is the parent: Now waiting for child with PID# 6
This is the parent: Child# 6 has exited
This is the parent: Now waiting for child with PID# 4
This is the parent: Child# 4 has exited
This is the parent: Now waiting for child with PID# 8
This is the parent: Child# 8 has exited
Turnaround (elapsed) time: 35 ticks.
Wait time: 7 ticks.
$
```

Test, instance 1

```

$ lab2 2

This program tests the correctness of your lab#2

Part c) testing that wait and waitpid interact correctly with the roundrobin scheduler:

The is child with PID# 11
Turnaround (elapsed) time: 1 ticks.
Wait time: 0 ticks.

The is child with PID# 12
Turnaround (elapsed) time: 1 ticks.
Wait time: 1 ticks.

The is child with PID# 13

The is child with PID# 14
Turnaround (elapsed) time: 2 ticks.
Wait time: 1 ticks.
The is child with PID# 10
Turnaround (elapsed) time: 2 ticks.
Wait time: 2 ticks.
Turnaround (elapsed) time: 2 ticks.
Wait time: 2 ticks.

This is the parent: Now waiting for child with PID# 13

This is the partent: Child# 13 has exited

This is the parent: Now waiting for child with PID# 11

This is the partent: Child# 11 has exited

This is the parent: Now waiting for child with PID# 12

This is the partent: Child# 12 has exited

This is the parent: Now waiting for child with PID# 10

This is the partent: Child# 10 has exited

This is the parent: Now waiting for child with PID# 14

This is the partent: Child# 14 has exited
Turnaround (elapsed) time: 32 ticks.
Wait time: 5 ticks.
$ █

```

Test, instance 2

The two above screenshots demonstrate the performance-tracking capabilities of my lab as well as the functionality of the round-robin scheduler. Not only does each processes get a turn, it doesn't favor any process specifically. Note that the **parent** processes necessarily take longer than their children because **they are using waitPID on them**. This was to verify that round-robin scheduling plays nicely with my wait, waitPID, and exit implementations.